

SMART TRAVEL ADVISOR

**A project report submitted in partial fulfillment of the requirements for the award of
credits to**

Machine Learning Course of

Bachelor of Technology

In

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING – ARTIFICIAL
INTELLIGENCE & MACHINE LEARNING (CSM)

By

B. Meghana

23BQ1A4218



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING – ARTIFICIAL INTELLIGENCE &
MACHINE LEARNING (CSM)

VASIREDDY VENKATADRI INSTITUTE OF TECHNOLOGY

(Approved by AICTE and permanently affiliated to JNTUK)

Accredited by NBA and NAAC with 'A' Grade

NAMBUR (V), PEDAKAKANI (M), GUNTUR-522 508

JUNE 2025

**DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING –
ARTIFICIAL INTELLIGENCE & MACHINE LEARNING (CSM)**

**VASIREDDY VENKATADRI INSTITUTE OF TECHNOLOGY:
NAMBUR**

**JAWAHARLAL NEHRU TECHNOLOGICAL UNIVERSITY
KAKINADA**



CERTIFICATE

This is to certify that the project titled “Smart travel advisor” is a Bonafide record of work done by Ms. B. Meghana under the guidance of Mrs. B. Lalitha Rajeswari, M. Tech, (Ph. D), Asst. Professor in partial fulfillment of the requirement for the award of credits to Machine Learning of Bachelor of Technology in Computer Science & Engineering – Artificial Intelligence & Machine Learning (CSM), JNTUK during the academic year 2024-25.

B. Lalitha Rajeswari

Course Instructor

Prof. K. Suresh Babu

Head of the Department

DECLARATION

I, B. MEGHANA (23BQ1A4218), hereby declare that the Project Report entitled “SMART TRAVEL ADVISOR” done by me under the guidance of Mrs. B. Lalitha Rajeswari, M.Tech,(Ph.D), Asst. Professor is submitted in partial fulfillment of the requirements for the award of degree of BACHELOR OF TECHNOLOGY in COMPUTER SCIENCE & ENGINEERING – ARTIFICIAL INTELLIGENCE & MACHINE LEARNING (CSM).

DATE :

SIGNATURE OF THE CANDIDATE

PLACE :

ACKNOWLEDGEMENT

We express our sincere thanks wherever it is due

We express our sincere thanks to the Chairman, Vasireddy Venkatadri Institute of Technology, Sri Vasireddy Vidya Sagar for providing us well equipped infrastructure and environment.

We thank Dr. Y. Mallikarjuna Reddy, Principal, Vasireddy Venkatadri Institute of Technology, Nambur, for providing us the resources for carrying out the project.

We express our sincere thanks to Dr. K. Giribabu, Dean of Academics for providing support and stimulating environment for developing the project.

Our sincere thanks to Dr. K. Suresh Babu, Head of the Department, Department of CSM, for his co-operation and guidance which helped us to make our project successful and complete in all aspects.

We also express our sincere thanks and are grateful to our guide Mrs. B. Lalitha Rajeswari , Associate Professor, Department of CSM, for motivating us to make our project successful and fully complete. We are grateful for his precious guidance and suggestions.

We also place our floral gratitude to all other teaching staff and lab technicians for their constant support and advice throughout the project.

NAME OF THE CANDIDATE

B. Meghana
(23BQ1A4222)

Table of Contents

CHAPTER 1: INTRODUCTION TO THE PROJECT	1
Aim of the project:	1
The Need for Personalized Nutrition	1
Objectives of the System	1
Why Use K-Means Clustering?	2
K-Means clustering is a powerful unsupervised machine learning algorithm used to group similar data points based on feature similarity. In the context of this project, each travel destination can be represented as a feature vector that includes:	2
□ Environmental attributes (e.g., temperature, humidity, air quality) ..	2
□ Cost indicators (e.g., accommodation, transportation, cost of living index)	2
□ Health-related metrics (e.g., regional disease risk, weather safety)	2
By applying K-Means, the system can create meaningful clusters of destinations such as “budget-friendly with moderate weather” or “luxury travel with high temperature and humidity,” allowing users to explore regions that align with their specific travel needs and constraints.	2
1.2 Problem Definition	2
The Modern Challenge of travel planning	2
In today’s dynamic lifestyle, travelers face increasing challenges in making decisions that align with both their health and financial well-being. Traditional trip planning methods—manual research, generalized travel blogs, and non-personalized travel apps—often fall short when it comes to adapting recommendations to individual health, preferences, and safety concerns.	2
Many people find it difficult to identify suitable destinations when factors like temperature sensitivity, health conditions, environmental hazards, and budget constraints are involved. Existing travel tools often lack personalization, making travel planning stressful and suboptimal.	2
<i>Complexity of Personalized Nutrition</i>	2

Information Overload and Decision Fatigue.....	3
Lack of Intelligent Personalization in Existing Systems	3
The Role of Clustering in Solving These Issues	4
To overcome these challenges, the Smart Travel Advisor implements K-Means clustering, an unsupervised learning algorithm that groups destinations based on a variety of relevant attributes, including:	4
❑ Weather conditions (e.g., temperature, humidity, air quality)	4
❑ Health factors (e.g., disease outbreak risk, altitude, allergen exposure).....	4
❑ Cost indicators (e.g., accommodation, food, transport).....	4
❑ Activity categories and local accessibility.....	4
Clustering solves multiple real-world problems:.....	4
❑ Reduces decision fatigue by presenting grouped destination options .	4
❑ Supports discovery of comparable destinations that meet specific needs	4
❑ Enables intelligent recommendations based on health and budgetary alignment rather than generic preferences	4
For example, a user seeking low-cost, asthma-friendly, moderate climate destinations will be shown clusters containing locations that meet all criteria—even if the user never searched for them by name.....	4
By applying clustering, the system introduces personalized, adaptive, and health-aware travel planning, empowering users to make informed and optimized travel decisions effortlessly.	4
1.3 What is Your Input.....	4
introduction to Input Requirements	4
User Input	4
Dataset Input	6
Each travel record or destination in the dataset typically contains the following fields:	6
❑ Destination Name.....	6
❑ Geographical Details	6
o Country, region, city.....	6

o	GPS coordinates.....	6
□	Climate Information.....	6
o	Average monthly temperature	6
o	Humidity levels.....	6
o	Air Quality Index (AQI).....	6
o	Rainfall and seasonal weather patterns.....	6
□	Health and Safety Indicators	6
o	Disease outbreak zones (e.g., malaria, dengue)	6
o	Altitude level and impact risk.....	6
o	Local medical facility availability	6
o	Travel advisories and crime/safety index	6
□	Cost Metrics	6
o	Average accommodation rates (budget to luxury)	6
o	Daily food and transportation costs.....	6
o	Activity-specific expenses.....	6
	1.4 What Output You Are Expecting.....	7
	Overview of Expected Outputs.....	7
	1. Destination Recommendation Output	7
	A ranked list of travel destinations tailored to individual user input, such as climate preference, health conditions, activity interest, budget, and seasonal availability.	7
	Key Characteristics:	7
□	Personalized to User Profile: Each user receives unique travel suggestions based on health needs, weather tolerance, or lifestyle preferences.....	7
□	Health and Safety Aligned: Recommendations account for medical conditions (e.g., altitude sensitivity, respiratory risks) and advisories.....	7
□	Diverse and Culturally Enriched: Suggested destinations include a variety of cultural experiences and activities aligned with the user’s interests.	7

□ Time-Sensitive Filtering: Users can filter suggestions based on available travel months or peak season suitability.	7
□ Relevance-Based Ranking: Destinations are ranked based on alignment with user constraints and goals.	7
1.5 What Are the Possible Steps Involved.....	9
Building a data-driven travel-recommendation and suitability-prediction platform calls for a systematic, phased approach. Each stage has specific technical processes, tools, and deliverables that collectively ensure the system is accurate, scalable, and genuinely helpful to travelers.	9
Major Phases	9
1. Requirement Analysis.....	9
2. Dataset Collection and Acquisition	9
3. Data Pre-processing and Cleaning	9
4. Feature Engineering and Selection	9
5. Clustering with K-Means	9
6. Recommendation Engine Design.....	9
7. Suitability & Cost-Prediction Models.....	9
8. User Interface (UI/UX) Development	9
9. Integration and Deployment.....	9
10. Testing and Evaluation	9
11. Feedback and Iteration	9
Each step underpins the next, ultimately delivering personalized, health-aware, and budget-conscious travel recommendations.....	9
A robust, multi-source dataset is the system’s backbone.	10
□ Open Data Portals & APIs	10
o World Bank (cost-of-living & safety indices)	10
o OpenWeather / Meteostat (historical climate)	10
o WHO / CDC health-risk bulletins	10
o Numbeo for crowd-sourced cost and pollution metrics	10
o UNWTO tourism statistics.....	10

□ Custom Web-Scraping / API Scripts to ingest accommodation prices, AQI, and seasonal trends.	10
□ Storage: centralized data lake (CSV/Parquet) or relational DB (PostgreSQL).	10
Data Preprocessing and Cleaning.....	10
1. Handle missing climate or cost values (mean/median imputation, region-level fallback).....	10
2. Normalize numeric features (temperature, AQI, hotel prices). ..	10
3. Encode categorical variables (continent, activity tags).....	10
4. Harmonize temporal granularity (convert all climate stats to monthly averages).	10
Feature Engineering and Selection.....	10
1. Composite Indices:.....	10
□ Comfort Index = temperature deviation + humidity factor	10
□ Affordability Score = (transport + accommodation + food) / PPP ..	10
2. One-Hot Tags: beach, mountain, heritage, festival, adventure. ..	10
3. Temporal Features: best-month flags, rainy-season indicator. ...	10
4. Health Flags: altitude > 2 500 m, dengue-risk zone, pollen-peak months.	10
Recommendation Engine Design.....	10
Suitability & Cost-Prediction Models	11
Integration and Deployment	11
Containerize services with Docker.....	11
Deploy on cloud (AWS/GCP/Azure) with CI/CD pipeline.....	11
Use a geospatial database extension (PostGIS) for efficient map queries.	11
Testing and Evaluation	11
1.6 How we obtained the Input Data / Existing Database.....	12
Introduction	12
A data-driven project such as Smart Travel Advisor is only as good as the quality, completeness, and relevance of its data. In our case, the	

input files centre on destinations, climate statistics, cost indices, health & safety advisories, and user context. These datasets feed the machine-learning pipeline for clustering, suitability prediction, and personalised recommendations.....	12
This section details the sources, formats, contents, and structure of the datasets, how they were acquired, and why they were selected.	12
Types of Input Files Used	12
□ Destination Master Dataset: Contains city/region names with unique IDs, geo-coordinates, and country codes; serves as the core reference for destinations.	12
□ Climate & Environmental Dataset: Provides monthly data on temperature, humidity, rainfall, and air quality (AQI) to assess environmental conditions.	12
□ Cost-of-Living & Tourism Spend Dataset: Includes data on accommodation, food, transportation, and PPP-adjusted cost indices for budgeting.....	12
□ Health & Safety Dataset: Offers details on disease risks, crime/safety scores, altitude levels, and medical facilities to ensure traveler well-being.....	12
□ User Profile Data (Optional): Captures user preferences, travel dates, health conditions, and budget to enable personalized travel suggestions.	12
Sources of Datasets.....	12
Format of Input Files.....	12
Preprocessing and Cleaning of Input Files.....	12
Dataset Size and Coverage	13
Tools Used for Handling the Input Files	13
1.7 Scope of the Project.....	14
Introduction	14
Technical Scope	15
Strategic Scope	16
Target Audience.....	17

Limitations and Constraints	17
Some challenges and limitations include:	17
□ Data Gaps: Not all destinations have complete climate, health, and cost data	17
□ Model Assumptions: K-Means clustering may not represent all types of destination groupings optimally	17
□ Personalization Dependency: More accurate recommendations require richer user profiles	17
□ Not Real-Time: Most climate and cost data is static or historical, not live	17
Chapter 2: Literature Review	22
2.1 Existing Methods and Algorithms	22
Chapter – 3. Experimental or materials and methods	25
Chapter 4: Results and Discussion, Performance Analysis	30
4.2 Optimal Number of Clusters.....	30
Using the Elbow Method on the Within-Cluster Sum of Squares (WCSS) curve for $k = 1 \dots 10$, the inflection clearly appeared at $k = 4$. This choice balanced cluster compactness with human interpretability.	30
4.3 Cluster Distribution and Centroid Highlights	30
4.4 Interpretation of Clusters.....	31
□ Cluster A suits health-sensitive travelers or families wanting moderate temperatures and very low pollution.	31
□ Cluster B appeals to budget backpackers who tolerate heat in exchange for affordable living costs.	31
□ Cluster C targets trekkers and adventure seekers comfortable with cooler weather and high elevation.	31
□ Cluster D attracts business or luxury tourists prepared to pay premium prices and accept denser urban environments	31
4.5 Visual Validation	32
Several plots confirmed the quality of grouping:	32

□ Temperature vs. AQI scatter colored by cluster showed distinct, minimally overlapping clouds.....	32
□ Cost-of-living histograms revealed non-intersecting peaks corresponding to each group.....	32
□ 3-D plots (Temperature, AQI, Cost) further illustrated clear separation, validating K-Means effectiveness.	32
4.6 Performance Analysis	32
□ K-Means trained on 8 k+ rows in under two seconds on a standard laptop, confirming feasibility for much larger or near-real-time datasets.....	33
4.7 Clustering driven Use-Case Example	33
□ Cluster A: Ideal for asthma patients or parents with infants who require low AQI and mild weather year-round.	33
□ Cluster B: Perfect for college students on shoestring budgets seeking warm beaches and street-food culture.	33
□ Cluster C: Suits trekkers, mountain cyclists, and digital detox retreats far from pollution.....	33
□ Cluster D: Fits executives on business trips or luxury travelers interested in fine dining and cosmopolitan nightlife.	33
Chapter 5: Summary, Conclusions, and Future Scope	34
<i>Summary</i>	34
Key steps completed were:	34
□ Collecting and merging data from weather APIs, air-quality feeds, cost-of-living portals, and health-safety bulletins to create a rich dataset of more than 8 000 city-month records.	34
□ Pre-processing that included unit standardisation (°F → °C, local currency → USD PPP), missing-value imputation, outlier capping, one-hot encoding of categorical flags, and Z-score normalization of numeric attributes.....	34
□ Implementing K-Means clustering to segment destinations by shared environmental and economic traits.	34
□ Selecting the optimal number of clusters (k = 4) using the Elbow Method, which balanced cluster compactness with interpretability.	34

□ Visualizing and interpreting results through scatter plots and heatmaps, confirming that clusters aligned with intuitive travel archetypes like “clean-air wellness towns” or “budget tropical hubs.” ..	34
□ Generating actionable recommendations that match user profiles to cluster centroids, thereby offering practical destination short-lists that respect health and budget constraints.	34
Through this pipeline, the project validated that unsupervised learning can meaningfully organize global travel data and deliver tailored, health-aware travel advice.	34
<i>Conclusions</i>	34
<i>Future Scope</i>	35
References	37
Journals.....	37
Books	37

ABSTRACT

This project presents a machine learning-based **Smart Travel Advisor** system that integrates travel recommendations with health and cost-aware predictions, focusing specifically on clustering and regression techniques. Utilizing diverse datasets encompassing trip details, weather conditions, cost indices, and health-related statistics, the system aims to assist users in selecting travel destinations aligned with their preferences, safety requirements, and budget constraints.

The core technique employed is **K-Means clustering**, which is used to group destinations based on shared characteristics such as temperature, humidity, and cost of living. These clusters represent distinct categories (e.g., budget-friendly, health-safe, luxury, weather-optimal) and form the basis for recommending travel locations tailored to the user's interests or constraints.

In addition to clustering, the system includes a **trip cost and health risk prediction module**, where statistical regression techniques (e.g., linear regression) are used to estimate total trip expenses and potential health risks based on features like duration, age of traveler, weather conditions, and environmental factors such as precipitation and leaf area index.

Overall, this project offers a scalable and interpretable solution for personalized travel planning, leveraging **unsupervised learning for clustering** and **supervised learning for prediction**. It showcases how machine learning can enable smarter, safer, and more budget-conscious travel decisions, promoting a better travel experience for users.

CHAPTER 1: INTRODUCTION TO THE PROJECT

Aim of the project:

In an age where health, cost-efficiency, and personalized experiences are gaining increasing importance, travelers seek more than just popular tourist spots—they desire destinations that align with their health conditions, weather preferences, and budget constraints. However, most existing travel platforms fall short of offering personalized travel recommendations that integrate environmental, health, and financial data.

The **Smart Travel Advisor** project aims to bridge this gap by developing an intelligent machine learning-based system that not only recommends travel destinations but also predicts trip costs and assesses travel suitability based on environmental and health-related factors. By leveraging **K-Means clustering**, the system categorizes destinations with similar weather and cost profiles, helping users explore options best suited to their preferences and health needs.

The core goal is to offer users a personalized and health-aware travel planning experience that supports smarter choices—whether it's finding budget-friendly trips with safe weather or planning around chronic health conditions like asthma or heat sensitivity.

The Need for Personalized Nutrition

Generic travel advice and one-size-fits-all itineraries often fail to meet individual needs. Factors such as temperature sensitivity, disease outbreaks, age, physical health, and travel budget significantly influence travel decisions. Although many websites offer destination suggestions, few provide **intelligent, personalized, and data-driven insights** based on health safety, cost, and environmental conditions.

This project addresses these limitations by introducing a **customizable, ML-powered travel advisor** that enables users to receive recommendations based on clusters of destinations matching their weather preferences, cost comfort, and safety levels. It empowers users to make informed travel decisions that are not only enjoyable but also responsible and well-aligned with personal well-being.

Objectives of the System

The **Smart Travel Advisor** system has been designed with the following core objectives:

- **Cluster Destinations Based on Environmental and Economic Data:** Use unsupervised learning to group locations based on weather patterns, humidity, and cost of living, helping users explore destinations that suit their climate preferences and budget.
- **Generate Personalized Travel Recommendations:** Consider user-specific factors like age, trip duration, and safety requirements (e.g., optimal weather range for health conditions) to suggest suitable destinations.
- **Predict Trip Costs and Health Risk Factors Accurately:** Use regression models to estimate total trip costs and potential health risks based on travel data, weather, and environmental features like temperature, precipitation, and air quality.

- **□Predict Trip Costs Accurately:** Estimate total travel expenses based on destination, accommodation, transportation, and user trip preferences, providing transparent cost breakdowns to aid in financial planning.
- **Provide an Easy-to-Use Interface:** Ensure that users can interact with the system effortlessly, without requiring technical expertise. Visual insights, intuitive inputs, and clear recommendations make the system accessible to all types of users.
- **Enable Learning from User Feedback:** Optionally, the system can adapt over time by learning from users' travel choices and preferences to improve the relevance and personalization of future recommendations.

Why Use K-Means Clustering?

K-Means clustering is a powerful unsupervised machine learning algorithm used to group similar data points based on feature similarity. In the context of this project, each travel destination can be represented as a feature vector that includes:

- **Environmental attributes** (e.g., temperature, humidity, air quality)
- **Cost indicators** (e.g., accommodation, transportation, cost of living index)
- **Health-related metrics** (e.g., regional disease risk, weather safety)

By applying K-Means, the system can create meaningful clusters of destinations such as “budget-friendly with moderate weather” or “luxury travel with high temperature and humidity,” allowing users to explore regions that align with their specific travel needs and constraints.

1.2 Problem Definition

The Modern Challenge of travel planning

In today's dynamic lifestyle, travelers face increasing challenges in making decisions that align with both their health and financial well-being. Traditional trip planning methods—manual research, generalized travel blogs, and non-personalized travel apps—often fall short when it comes to adapting recommendations to individual health, preferences, and safety concerns.

Many people find it difficult to identify suitable destinations when factors like temperature sensitivity, health conditions, environmental hazards, and budget constraints are involved. Existing travel tools often lack personalization, making travel planning stressful and suboptimal.

Complexity of Personalized Nutrition

- Travel suitability is not one-size-fits-all. Each traveler has a unique profile shaped by age, health conditions, climate preferences, and financial capacity. For example:
- A person with respiratory issues must avoid highly humid or polluted regions.
- A student with a tight budget needs low-cost travel options with safe accommodation.
- A senior citizen may prefer mild temperatures and minimal travel hassle.
- A family traveling with children might seek destinations with low health risks and easy accessibility.
- The **Smart Travel Advisor** addresses these diverse needs by offering a machine learning-driven platform that **recommends, predicts, and personalizes** travel plans

based on real-world data—making the travel experience not only enjoyable but also responsible and well-informed.

Unfortunately, current travel platforms fail to address the intersection of **personal preferences, health constraints, budget limitations, and environmental safety** in a unified and intelligent manner.

While some travel websites provide broad filters (e.g., price range, activity type), they **rarely integrate real-time health, climate, or personalized wellness data**—key factors that significantly influence a traveller’s decision-making. Travelers with specific needs (e.g., asthma, budget travel, or extreme weather intolerance) often lack access to tools that consider **all factors simultaneously** when recommending destinations.

Information Overload and Decision Fatigue

Another core challenge is **decision fatigue**. Users are bombarded with thousands of destinations, deals, and itineraries, yet **lack intelligent filters** that prioritize what truly matters to them:

- *“What destinations align with my health conditions this season?”*
- *“Where can I travel safely within my budget this month?”*
- *“Which places offer mild weather, low air pollution, and local transport options?”*

Without context-aware filtering, users feel overwhelmed—leading to poor travel decisions, suboptimal experiences, or repetitive plans that don’t match their evolving goals.

Lack of Intelligent Personalization in Existing Systems

Most mainstream travel platforms rely on **static filtering or keyword searches** (like “beach,” “budget,” or “romantic getaway”). These approaches:

- **Fail to account for individual health or weather sensitivities**
- **Don’t adapt to past travel behavior or preferences**
- **Cannot suggest alternatives that offer similar experiences in safer or more suitable conditions**

Moreover, **machine learning-based clustering or classification** is almost absent in existing travel recommendation systems. As a result, users miss out on discovering hidden patterns—such as destinations with similar **climate, cost range, safety levels, or health conditions compatibility**.

The Role of Clustering in Solving These Issues

To overcome these challenges, the **Smart Travel Advisor** implements **K-Means clustering**, an unsupervised learning algorithm that groups destinations based on a variety of relevant attributes, including:

- Weather conditions (e.g., temperature, humidity, air quality)
- Health factors (e.g., disease outbreak risk, altitude, allergen exposure)
- Cost indicators (e.g., accommodation, food, transport)
- Activity categories and local accessibility

Clustering solves multiple real-world problems:

- **Reduces decision fatigue** by presenting grouped destination options
- **Supports discovery** of comparable destinations that meet specific needs
- **Enables intelligent recommendations** based on health and budgetary alignment rather than generic preferences

For example, a user seeking **low-cost, asthma-friendly, moderate climate destinations** will be shown clusters containing locations that meet all criteria—even if the user never searched for them by name.

By applying clustering, the system introduces **personalized, adaptive, and health-aware travel planning**, empowering users to make **informed and optimized travel decisions** effortlessly.

1.3 What is Your Input

introduction to Input Requirements

To determine the **accuracy and efficiency** of recommendations and predictions, input plays a critical role in the Smart Travel Advisor system. This project relies on two primary categories of input:

1. **User Input** – Information provided by users about their travel preferences, health conditions, and personal constraints.
2. **Dataset Input** – External data sources, such as historical travel trends, climate records, air quality indices, safety ratings, and cost estimations used to train and power the recommendation engine.

In any machine learning or data-driven system, the **quality, consistency, and structure** of input data are vital. They ensure the system delivers **meaningful, context-aware travel suggestions** and accurate estimations of travel conditions.

User Input

User input is a **dynamic and personalized** component of the system. It captures individual travel needs and constraints based on factors such as **health sensitivity, budget, purpose of travel, and seasonal preferences**.

These inputs guide the system in tailoring destination recommendations that are not only enjoyable but also **safe, affordable, and well-suited** to the user's conditions.

Key Components of User Input

- **Travel Purpose**
 - Leisure
 - Adventure
 - Wellness/Retreat
 - Business
 - Family Vacation
- **Health Conditions**
 - Respiratory issues (e.g., asthma)
 - Allergy sensitivity (pollen, dust, food, etc)
 - Altitude sensitivity
 - Immuno-compromised conditions
- **Budget Constraints**
 - Daily budget (low, medium, high)
 - Accommodation preference (hostel, hotel, resort, Airbnb)
 - Transportation choices (public, private, walkable cities)
- **Climate Preferences**
 - Cold
 - Tropical
 - Moderate
 - Dry/Arid
- **Time Constraints**
 - Travel duration (e.g., weekend gateway, 1-week trip)
 - Preferred months or seasons
- **Activity Interests**
 - Beaches
 - Mountains
 - Historical sites
 - Wildlife
 - Urban exploration
- **Travel Group Type**
 - Solo
 - Couple
 - Friends
 - Family with kids
 - Seniors

All data is cleaned, normalized, and integrated into feature vectors to enable efficient clustering using machine learning algorithms like **K-Means**.

Dataset Input

The **dataset input** forms the foundation of the **Smart Travel Advisor's** recommendation and prediction engine. This data comprises thousands of destination records and travel insights, each annotated with detailed information such as environmental factors, seasonal trends, cost metrics, health advisories, and user reviews.

This structured and diverse dataset is essential for training machine learning models that can intelligently cluster destinations, predict optimal travel times, and generate personalized travel recommendations based on user profiles.

Core Attributes in the Travel Dataset

Each travel record or destination in the dataset typically contains the following fields:

- **Destination Name**
- **Geographical Details**
 - Country, region, city
 - GPS coordinates
- **Climate Information**
 - Average monthly temperature
 - Humidity levels
 - Air Quality Index (AQI)
 - Rainfall and seasonal weather patterns
- **Health and Safety Indicators**
 - Disease outbreak zones (e.g., malaria, dengue)
 - Altitude level and impact risk
 - Local medical facility availability
 - Travel advisories and crime/safety index
- **Cost Metrics**
 - Average accommodation rates (budget to luxury)
 - Daily food and transportation costs
 - Activity-specific expenses

1.4 What Output You Are Expecting

Overview of Expected Outputs

The primary goal of the **Smart Travel Advisor** project is to provide users with **personalized travel destination recommendations** based on their preferences and to deliver **accurate predictions** related to travel suitability—such as climate conditions, health advisories, cost estimates, and activity options. The system is designed to produce both **user-centric outputs** (the interactive or visible parts) and **model-centric outputs** (machine learning-driven internal decisions).

The expected outputs are organized into the following categories:

1. Destination Recommendation Output

A ranked list of travel destinations tailored to individual user input, such as climate preference, health conditions, activity interest, budget, and seasonal availability.

Key Characteristics:

- **Personalized to User Profile:** Each user receives unique travel suggestions based on health needs, weather tolerance, or lifestyle preferences.
- **Health and Safety Aligned:** Recommendations account for medical conditions (e.g., altitude sensitivity, respiratory risks) and advisories.
- **Diverse and Culturally Enriched:** Suggested destinations include a variety of cultural experiences and activities aligned with the user's interests.
- **Time-Sensitive Filtering:** Users can filter suggestions based on available travel months or peak season suitability.
- **Relevance-Based Ranking:** Destinations are ranked based on alignment with user constraints and goals.

2. Clustering Output (K-Means Results)

Using K-Means clustering, the system groups destinations with similar environmental, economic, and experiential profiles. This supports both backend structuring and intelligent user-facing recommendations.

Clustering Output Features:

- **Cluster Centroids:** Each centroid defines the average conditions (e.g., temperature, cost, air quality) of destinations within a cluster.
- **Cluster Labels:** Destinations are tagged with thematic labels such as:
 - Budget-Friendly Coastal Getaways
 - High-Altitude Adventure Spots
 - Mild-Climate Wellness Retreats

- Similarity Insights: Alternative destinations within the same cluster are suggested, enhancing variety while maintaining relevance.

3. Suitability and Risk Prediction Output

The system offers predictive insights such as expected weather, potential health risks, and overall travel suitability for each location based on season and user profile.

Predicted Outputs Include:

- Average temperature and humidity
- Pollution levels (AQI)
- Disease risk zones
- Safety score predictions
- Cost estimate projections

4. User Interface Output

A clean, responsive interface where users can interact with results—view recommendations, explore cluster insights, and receive guidance.

UI Highlights:

- Interactive maps and weather overlays
- Filtered destination cards based on user constraints
- Nutritional and health indicators for destinations (e.g., clean food availability, safe water, etc.)

5. System Logs and Metrics

Backend outputs used for monitoring, system performance tracking, and continuous model improvement.

Examples:

- User interaction logs for preference learning
- Recommendation accuracy and feedback collection
- Cluster quality metrics (e.g., Silhouette Score)

1.5 What Are the Possible Steps Involved

Building a data-driven **travel-recommendation and suitability-prediction platform** calls for a systematic, phased approach. Each stage has specific technical processes, tools, and deliverables that collectively ensure the system is accurate, scalable, and genuinely helpful to travelers.

Major Phases

1. Requirement Analysis
2. Dataset Collection and Acquisition
3. Data Pre-processing and Cleaning
4. Feature Engineering and Selection
5. Clustering with K-Means
6. Recommendation Engine Design
7. Suitability & Cost-Prediction Models
8. User Interface (UI/UX) Development
9. Integration and Deployment
10. Testing and Evaluation
11. Feedback and Iteration

Each step underpins the next, ultimately delivering personalized, health-aware, and budget-conscious travel recommendations.

Requirement analysis

Target Users: budget travelers, families, adventure seekers, health-sensitive tourists, business travelers.

Key Features:

- Destination filtering by climate, health risk, and cost
- Real-time suitability scores and cost estimates
- Interactive maps and cluster-based discovery

Project Objectives: personalization, safety, cost transparency, ease-of-use.

Tech Stack: Python, Pandas, Scikit-learn, GeoPandas, Flask / Django (backend), PostgreSQL, HTML/CSS/JS (frontend), Map libraries (Leaflet or Mapbox).

Dataset Collection and Acquisition

A robust, multi-source dataset is the system's backbone.

- Open Data Portals & APIs
 - World Bank (cost-of-living & safety indices)
 - OpenWeather / Meteostat (historical climate)
 - WHO / CDC health-risk bulletins
 - Numbeo for crowd-sourced cost and pollution metrics
 - UNWTO tourism statistics
- **Custom Web-Scraping / API Scripts** to ingest accommodation prices, AQI, and seasonal trends.
- **Storage:** centralized data lake (CSV/Parquet) or relational DB (PostgreSQL).

Data Preprocessing and Cleaning

1. Handle missing climate or cost values (mean/median imputation, region-level fallback).
2. Normalize numeric features (temperature, AQI, hotel prices).
3. Encode categorical variables (continent, activity tags).
4. Harmonize temporal granularity (convert all climate stats to monthly averages).

Feature Engineering and Selection

1. Composite Indices:
 - Comfort Index = temperature deviation + humidity factor
 - Affordability Score = (transport + accommodation + food) / PPP
2. **One-Hot Tags:** beach, mountain, heritage, festival, adventure.
3. **Temporal Features:** best-month flags, rainy-season indicator.
4. **Health Flags:** altitude > 2 500 m, dengue-risk zone, pollen-peak months.

Clustering with K-Means

1. Input: standardized vectors [Comfort Index, Affordability, AQI, Health Risk, Activity Tags].
2. Optimal k chosen via Elbow / Silhouette.
3. Output clusters such as:
 - “Mild-climate, mid-budget coastal towns”
 - “High-altitude adventure hubs”
 - “Low-cost cultural cities with moderate AQI”.

Recommendation Engine Design

- **Filtering:** Hard constraints (budget cap, health exclusions).

- **Ranking:** distance to cluster centroid + personalized weights (climate vs. cost vs. safety).
- **Diversity Boost:** penalize over-similar suggestions to avoid monotony.

Suitability & Cost-Prediction Models

Regression Models (Linear / Gradient Boosting) for total-trip cost based on stay length, season, and accommodation class.

Classification Models (Random Forest, XGBoost) for health suitability (Safe / Caution / Avoid) using climate and disease indicators.

User Interface (UI/UX) Development

Frontend: responsive dashboard with filters, map visualization, and card-style destination summaries.

APIs: REST or GraphQL endpoints serving recommendations, cluster info, and predictions in JSON.

Accessibility: color-blind-friendly palettes, readable fonts, low-bandwidth mode.

Integration and Deployment

Containerize services with Docker.

Deploy on cloud (AWS/GCP/Azure) with CI/CD pipeline.

Use a geospatial database extension (PostGIS) for efficient map queries.

Testing and Evaluation

Unit & Integration Tests for data pipelines and API endpoints.

Model Metrics:

- Clustering quality (Silhouette Score)
- Cost-prediction RMSE
- Suitability-classification F1 Score

User Testing: A/B experiments on recommendation relevance.

Feedback and Iteration

Collect implicit feedback (click-through, destination saves).

Periodically retrain models with new data.

Introduce reinforcement-learning or bandit approaches for continuous personalization.

1.6 How we obtained the Input Data / Existing Database

Introduction

A data-driven project such as Smart Travel Advisor is only as good as the quality, completeness, and relevance of its data. In our case, the input files Centre on destinations, climate statistics, cost indices, health & safety advisories, and user context. These datasets feed the machine-learning pipeline for clustering, suitability prediction, and personalised recommendations.

This section details the **sources, formats, contents, and structure** of the datasets, how they were acquired, and why they were selected.

Types of Input Files Used

The input data used in this project includes:

- **Destination Master Dataset:** Contains city/region names with unique IDs, geo-coordinates, and country codes; serves as the core reference for destinations.
- **Climate & Environmental Dataset:** Provides monthly data on temperature, humidity, rainfall, and air quality (AQI) to assess environmental conditions.
- **Cost-of-Living & Tourism Spend Dataset:** Includes data on accommodation, food, transportation, and PPP-adjusted cost indices for budgeting.
- **Health & Safety Dataset:** Offers details on disease risks, crime/safety scores, altitude levels, and medical facilities to ensure traveler well-being.
- **User Profile Data (Optional):** Captures user preferences, travel dates, health conditions, and budget to enable personalized travel suggestions.

Sources of Datasets

The datasets used were obtained from **trusted, open-source, or publicly available platforms** that are frequently used for academic and commercial data science projects. The main sources include:

Format of Input Files

- **Destination Dataset:** City names, IDs, coordinates, and country codes.
- **Climate Dataset:** Monthly temperature, humidity, rainfall, and AQI.
- **Cost Dataset:** Accommodation, food, transport, and cost indices
- **User Data (Optional):** Preferences, health, budget, and travel dates.

Preprocessing and Cleaning of Input Files

Before using the data, several preprocessing tasks were conducted:

- **Unit Standardisation:** Convert °F to °C and local currencies to USD (PPP).
- **Temporal Alignment:** Convert raw weather data to monthly averages.
- **Missing Values:** Use regional interpolation and median imputation.
- **Outlier Treatment:** Winsorize extreme AQI and crime values.

These steps are essential to ensure uniformity and reliability of the data used for modeling.

Dataset Size and Coverage

- **Destinations:** $\approx$ 8 000 cities/regions across 195 countries
- **Climate Records:** > 1 000 000 monthly rows (10 yr window)
- **Cost Metrics:** > 25 cost attributes per destination
- **Health & Safety Flags:** 12 indicators (disease, altitude, crime, hospital density)
- **Activity Tags:** 15–20 per destination (beach, trekking, heritage, etc.)

This scale ensures robust clustering and fine-grained personalisation.

How the Input Files Are Utilized

- **K-Means Clustering** → feature vector = [Climate Comfort, AQI, Cost Index, Safety Score, Activity Tags]
- **Regression Models** → trip-cost prediction using stay length + season + accommodation type
- **Classification Models** → health suitability (Safe / Caution / Avoid) based on climate & disease data
- **Frontend Maps & Filters** → pull structured data for interactive exploration
- **Personalisation Layer** → filter clusters by user budget, health flags, preferred climates.

Tools Used for Handling the Input Files

- **Python (Pandas, NumPy, GeoPandas)** – ETL, cleaning, geospatial joins
- **Scikit-learn** – clustering and predictive modelling
- **PostgreSQL + PostGIS** – centralised storage & spatial queries
- **Jupyter Notebooks** – EDA and feature engineering
- **Airflow** – scheduled data refresh / API harvesting
- **Leaflet / Mapbox** – frontend map rendering

These tools guarantee that data flows cleanly from raw sources to production-ready features, underpinning the **Smart Travel Advisor**'s reliability and accuracy

1.7 Scope of the Project

Introduction

The scope of a project outlines its boundaries, capabilities, and intended impact. In the context of the *Smart Travel Advisor*, the scope goes beyond a basic trip recommender. It integrates unsupervised learning, environmental data, economic factors, and user preferences to deliver a personalized and holistic travel planning experience. This section describes the functional, technical, and strategic scope of the system.

Functional Scope

The functional scope refers to the specific tasks and features the system is expected to perform.

1. Personalized Destination Recommendation

- Recommend cities or regions based on:
 - Budget constraints
 - Health preferences (e.g., low AQI, altitude sensitivity)
- Use clustering to group destinations and match user input to ideal travel clusters.

2. Environmental and Cost-Based Filtering

- Enable filtering of destinations based on:
 - Weather conditions (e.g., temperature, humidity)
 - Air Quality Index (AQI)

3. Custom Travel Profiles

- Save user preferences like:
 - Preferred continent or region
 - Travel month or duration
 - Budget and accommodation style
 - Activity interest (e.g., hiking, beaches, heritage sites)

4. Activity And Safety Insights

- Highlight activities at each destination (e.g., festivals, sightseeing)
- Provide health and safety scores:
 - Disease risk levels
 - Crime/safety ratings
 - Proximity to medical facilities

5. Interactive and Intuitive User Interface

- Intuitive interface for:
 - Inputting travel preferences
 - Viewing destination profiles
 - Comparing recommended locations with graphs and scores
- Responsive design for both desktop and mobile platforms

Technical Scope

The technical scope focuses on the architecture, algorithms, and tools used to power the system:

1. Machine Learning and Clustering

- Apply K-Means (or other clustering algorithms) on destination features:
 - Climate
 - Cost
 - Safety
 - Environmental quality

2. Multi-Feature Travel Matching

- Implement ranking systems based on:
 - Similarity between user input and destination cluster centroids
 - Normalized scores (e.g., cost, health safety, climate)

3. Data Engineering and Preprocessing

- Handle and preprocess datasets involving:
 - Weather (monthly temperature, rainfall, humidity)
 - Safety and health indices
 - Currency conversions (USD PPP)
- Normalize, clean, and align datasets across time and location dimensions.

4. Frontend and Backend Development

- Full-Stack implementation using:
 - **Frontend:** HTML, CSS, JavaScript, or React
 - **Backend:** Python with Flask or Django
 - **Database:** SQLite or MySQL for storing user and recipe data
- Deploy on platforms like Heroku or AWS.

5. API Integration (Optional)

- Connect to APIs for:
 - Real-time weather and pollution updates
 - Tourist activity and event data
 - Country travel advisories or visa requirements

Strategic Scope

This defines the long-term vision and practical utility of the system:

1.Travel and Wellness Support

- Promote health-conscious travel by factoring air quality, disease prevalence, and altitude
- Help users avoid locations that may negatively impact chronic conditions

2.Budget-Aware Planning

- Suggest destinations within a user's financial means using PPP-adjusted cost indices
- Help optimize travel seasons to avoid peak prices

3.Educational Application

- Use as a teaching aid for:
 - Geography and Climate Education
 - Data science applications in tourism
 - Comparative cost-of-living analysis

4.System Expansion Opportunities

- Future integration with:
 - Fitness apps or wearables for syncing travel activity goals
 - Hotel/flight booking platforms
 - Local transportation and guide services

Out-of-Scope Elements

To maintain feasibility in the initial phase, the following features are **excluded**:

- Real-time voice assistants or virtual travel agents
- Hotel/flight booking functionalities
- Currency exchange predictions or visa processing

- Full integration with social media for travel sharing
- Dynamic route planning or offline maps

These may be considered in future updates.

Target Audience

This Smart Travel Advisor targets a broad demographic:

- Health-conscious travelers looking for safe and climate-suitable places
- Budget travelers wanting maximum value
- Students and researchers studying travel trends or data science
- Planners needing data-backed travel suggestions
- Families and solo travelers alike

Limitations and Constraints

Some challenges and limitations include:

- Data Gaps: Not all destinations have complete climate, health, and cost data
- Model Assumptions: K-Means clustering may not represent all types of destination groupings optimally
- Personalization Dependency: More accurate recommendations require richer user profiles
- Not Real-Time: Most climate and cost data is static or historical, not live

1.8 Applications

Introduction

The **Smart Travel Advisor** delivers value well beyond simple destination search. By merging machine-learning insights with environmental, economic, and health data, it supports decision-making across tourism, wellness, corporate planning, education, and public policy. Below are key domains where the system can be deployed to enhance user experience, optimize costs, and promote safer, more sustainable travel.

1. Personalized Trip-Planning Platforms

Use case: Mobile or web apps that craft door-to-door itineraries.

- Tailored destination short-lists based on budget, health, climate, and activity interests.
- Seasonality checks to avoid extreme weather or peak-price periods.
- Dynamic cost forecasts (flights + stay + food) plotted against user budget.

- One-click re-ranking when a traveler’s priorities (e.g., “cheapest” vs. “clean air”) change.

2. Health & Wellness-Focused Services

Use case: Platforms or clinics helping clients manage chronic conditions while traveling.

- Low-AQI or low-altitude suggestions for respiratory or cardiac patients.
- Heat-risk alerts for seniors or pregnant travelers.
- Nearest medical facility mapping in every recommended city.
- Integration with wearable data (heart rate, oxygen saturation) to refine future advice.

3. Online Travel Agencies (OTAs) & Meta-Search Engines

Use case: Expedia-, Skyscanner-, or booking-style portals.

- Context-aware filters that blend cost, climate, safety, and user health in a single slider interface.
- Cluster-based “Similar to” panels (“Liked Bali? Try Phu Quoc—same weather, lower price”).
- Real-time upsell of carbon-offset options or off-season discounts.

4. Corporate & Remote-Work Mobility Tools

Use case: HR or travel-expense platforms managing employee trips.

- Policy-compliant routing (budget caps, safety thresholds).
- Health-risk dashboards for duty-of-care teams.
- Cost-projection APIs feeding expense approvals and per-diem calculations.
- Well-being scores to encourage travel to low-stress, amenity-rich hubs for remote work.

5. Destination Marketing & Tourism Boards

Use case: Government or regional DMOs seeking data-driven promotion.

- Cluster analysis to identify undermarketed regions sharing traits with popular spots.
- Visitor archetype mapping to craft targeted campaigns (e.g., “budget-friendly wellness retreats”).
- Early warning on climate or health factors impacting tourism flows.

6. Insurance & Risk-Assessment Firms

Use case: Travel insurers or banks offering trip-protection products.

- Granular risk scores (disease, crime, weather volatility) feed premium calculators.
- Destination-specific advisories pushed policyholders in real time.
- Preventive alerts reducing claim frequency and severity.

7. Educational & Research Platforms

Use case: Universities and data-science programs.

- Hands-on datasets for courses in machine learning, geography, or environmental economics.
- Visualization tools teach clustering, regression, and anomaly detection.
- Case studies on sustainable tourism and climate resilience.

8. Smart-City & Public-Policy Planning

Use case: Urban planners forecasting tourist load and infrastructure needs.

- Visitor-flow predictions by season for transport and utilities planning.
- Impact modeling of climate change on tourism revenue.
- Scenario testing (e.g., AQI spikes, currency shocks) to guide resilience strategies.

Target Beneficiaries

- Individual travelers seeking safe, affordable, and health-compatible trips.
- Travel agencies & OTAs need richer personalization.
- Corporate mobility & HR teams optimizing cost and employee wellness.
- Healthcare providers advising patients on travel.
- Insurance, banking, and government sectors assessing destination risk.
- Academics & students exploring real-world ML applications in travel.

Limitations & Future Opportunities

- Data Gaps: Remote or emerging destinations may lack complete climate or cost data.
- Model Bias: Clustering may underrepresent niche travel interests (e.g., extreme sports).
- Live Data Dependency: Real-time health advisories and flight costs require continuous API integration.
- Expansion Paths: Incorporate booking engines, carbon-footprint calculators, or voice-assistant interfaces in later phases.

.

1.9 Use Cases

Use Case 1

Budget-Conscious Student – Affordable, Safe Destinations

User: Rahul, 20, undergraduate on a limited budget.

Scenario:

Rahul wants an affordable 5-day getaway during his semester break in July.

1. He logs in and selects “**budget travel**” as the priority, sets a **₹2 000/day** ceiling, and chooses **temperate climate** with **low AQI**.

2. The Smart Travel Advisor:
 - **Filters** destinations by cost-of-living indices and seasonal flight prices.
 - **Cluster's** locations with similar climate and safety scores.
 - **Ranks** options by total projected cost vs. budget, highlighting off-season discounts.

Outcome

Rahul receives a short-list of three destinations under budget, each with cost breakdowns, hostel suggestions, and an air-quality “green check,” enabling him to book confidently without hours of research.

Use Case 2

Asthma Patient – Health-Safe Trip Planning

User: Priya, 55, teacher with chronic asthma.

Scenario

Priya is planning a 10-day holiday in December but must avoid polluted or high-altitude regions.

1. She enters “**low AQI,**” “**altitude < 1 000 m,**” and “**medical facilities nearby**” as non-negotiable filters.
2. The system:
 - **Excludes** destinations with historical AQI > 75 or altitude > 1 000 m.
 - **Uses** clustering to surface coastal and forest locations with stable, clean air.
 - **Displays** nearest hospitals and real-time health advisories.

Outcome

Priya receives destination cards showing expected AQI, humidity, and hospital proximity, ensuring she can travel safely without exacerbating her asthma.

Use Case 3

Remote-Work Nomad – Climate & Connectivity Match

User: Alex, 32, software engineer working remotely.

Scenario:

Alex seeks a one-month stay where he can surf in the mornings and work reliably online.

1. He sets preferences for “**mild climate (20–28 °C), beach activities, high internet speed, mid-range cost.**”
2. The advisor:
 - **Clusters** coastal cities with similar temperatures and cost profiles.
 - **Scores** internet quality using crowd-sourced speed data.

- **Recommends** coworking spaces and monthly apartment rentals.

Outcome

Alex gets three “digital-nomad friendly” spots with predicted monthly expenses, average 4G/5G speeds, and nearby surf schools—allowing him to pick the best fit for work-life balance in minutes.

1.10 Conclusion

The primary goal of this project is to design and implement a **smart, data-driven travel advisory system** that recommends personalized travel destinations based on users’ preferences, health conditions, budgets, and environmental comfort. By leveraging techniques like **K-Means clustering**, data preprocessing, and intelligent filtering, the system identifies suitable travel options from large datasets, helping users make informed, safe, and optimized travel decisions.

In an era where **travel planning is overwhelmed with choices**, individuals often struggle to match destinations with their **personal constraints**, be it **budget, climate suitability, health considerations, or safety concerns**. Traditional travel platforms tend to focus on popularity or reviews, without accounting for personalized health risks, affordability, or environmental factors.

This project bridges that gap using **machine learning and real-world datasets** to automate destination discovery, health-aware travel recommendations, and location-based clustering. It empowers users with **data-backed insights** rather than guesswork, making travel **safer, smarter, and more accessible**, especially for those with specific health needs, budget limits, or lifestyle preferences.

Ultimately, the Smart Travel Advisor not only simplifies travel planning but also demonstrates the broader utility of **AI in solving human-centric problems** through personalization, prediction, and intelligent design.

Chapter 2: Literature Review

The development of a Smart Travel Advisor system is a multidisciplinary task involving **machine learning, data analytics, environmental science, and user experience design**. To build a robust and intelligent advisor, it is important to review existing techniques used in recommendation systems, clustering algorithms, and travel-related data modelling. This chapter outlines key approaches used in similar domains, assesses their advantages and limitations, and identifies essential libraries and tools that will power the implementation of the proposed system.

2.1 Existing Methods and Algorithms

1. Collaborative Filtering (CF)

Collaborative filtering is widely used in many recommendation systems. It makes predictions by learning from user interactions—if users with similar profiles visited similar destinations, the system recommends places based on collective behaviour. In travel applications, CF can suggest destinations that similar travelers enjoyed. However, CF heavily relies on historical data and user engagement. It suffers from the **cold start problem**, making it less effective for new users or less-visited locations. Moreover, it cannot incorporate weather suitability, health conditions, or real-time constraints.

2. Content-Based Filtering

Content-based filtering focuses on user preferences and item features such as destination climate, terrain, activities offered, and cost. For example, a user who prefers mountain trekking and cool climates will receive recommendations for hill stations with similar environmental parameters. This method works well for new users since it only needs input preferences. However, it tends to lack diversity and may repeatedly suggest similar locations unless carefully tuned.

3. Clustering Algorithms (K-Means, DBSCAN)

Clustering algorithms like **K-Means** are particularly useful in grouping destinations based on **climate conditions, affordability, air quality index (AQI), health comfort, and more**. For instance, using unsupervised learning, destinations can be grouped into clusters like "budget-friendly warm beach locations" or "cool, high-altitude destinations with low pollution." This helps users explore new destinations aligned with their comfort zone and constraints. In the project, **K-Means clustering** is implemented using Scikit-learn as follows:

4. Decision Trees & Random Forests

Supervised learning models like Decision Trees and Random Forests are valuable for **predicting travel suitability scores** based on labeled historical data. These models can be trained to classify whether a location is suitable for a user with specific health constraints

(e.g., asthma or heat intolerance). They provide **interpretable and robust predictions**, though they require well-labeled datasets to function effectively.

5. *Deep Learning Techniques*

Advanced systems may employ deep learning for applications such as **image-based destination recognition, natural language understanding of reviews, or sequence modeling for travel itineraries**. For example, CNNs can classify destination types based on images, while RNNs or Transformers can generate optimized multi-day travel routes. However, these models are data- and compute-intensive and may not be necessary in the initial system development.

2.2 Key Machine Learning Libraries and Tools

To implement the Smart Travel Advisor, several Python libraries form the backbone of data handling, modeling, and visualization.

1. *Scikit-learn (sklearn)*

Scikit-learn is a robust library offering clustering, classification, regression, and preprocessing tools. It is central to this project for implementing **K-Means clustering**, scaling data, and evaluating models, as demonstrated in the code snippet above.

2. *Pandas*

Pandas is essential for data ingestion, cleaning, and manipulation. It allows us to handle destination datasets, filter locations by features like region or cost, and compute summary statistics.

```
import pandas as pd

df = pd.read_csv('destinations.csv')
print(df[['PlaceName', 'Temperature', 'AQI']].head())
```

This enables easy inspection and filtering of destination attributes for modeling.

3. *NumPy*

NumPy is used for numerical operations and efficient matrix computations. For example, it helps calculate average environmental parameters across multiple destinations:

```
import numpy as np

climate_data = np.array([[30, 70], [20, 50], [35, 90]]) # [temp, AQI]
mean_values = np.mean(climate_data, axis=0)
print("Mean Climate Metrics:", mean_values)
```

4. Matplotlib and Seaborn

For visualizing destination clusters, weather trends, or health-risk zones, Matplotlib and Seaborn provide intuitive plots. The following snippet shows how to visualize clusters formed using K-Means:

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.scatterplot(x='Temperature', y='AQI', hue='Cluster', data=df)
plt.title('K-Means Clustering of Travel Destinations')
plt.xlabel('Temperature (°C)')
plt.ylabel('Air Quality Index')
plt.show()
```

These visualizations help users and developers better interpret the groupings and understand location-based patterns.

Chapter – 3. Experimental or materials and methods

3.1 Introduction

This chapter presents the methodology followed to develop the **Smart Travel Advisor** data-driven system that recommends travel destinations based on user preferences such as climate, air quality, cost, health safety, and seasonal patterns. It outlines the system components, tools used, dataset descriptions, preprocessing steps, clustering techniques, and the prediction methodology adopted.

3.2 System Overview

The Smart Travel Advisor functions as a hybrid system that combines both recommendation and prediction capabilities. Its main objectives are:

- To suggest travel destinations that align with the user's budget, climate preferences, air quality needs, and activity interests.
- To predict missing information such as estimated cost, air quality index, or climate comfort scores for destinations with incomplete data.
- To group destinations into clusters based on common characteristics using the K-Means algorithm. These clusters help in categorizing places such as “affordable coastal cities” or “clean-air hill stations,” making recommendations more meaningful and personalized.

3.3 Tools and Software Used

The development was done in Python, chosen for its simplicity and the vast ecosystem of libraries supporting machine learning and data analysis.

- Scikit-learn was used to implement K-Means clustering and regression models. It also provided tools for preprocessing like normalization and model evaluation.
- Pandas played a major role in reading, cleaning, and analyzing tabular data, especially datasets related to destination cost, climate, and health conditions.
- NumPy was used for efficient numerical computations such as scaling, distance calculations, and vector operations.
- Matplotlib and Seaborn were employed for visualizing destination clusters, temperature distributions, and air quality patterns.
- Jupyter Notebook served as the development environment for writing, testing, and documenting code and visual outputs.
- GeoPandas was optionally considered for matching cities to geographical climate zones.
- CSV files were used for storing datasets, and SQLite was considered when structured, persistent storage was needed.

3.4 Dataset Description

Multiple datasets were gathered from trusted public sources:

- OpenWeather and Meteostat APIs provided monthly temperature, rainfall, and humidity data for various cities.
- IQAir and WAQI offered air quality data, including seasonal AQI readings.
- Numbeo contributed cost-of-living data including accommodation, food, and local transportation costs.
- WHO and CDC data were used to flag destinations for disease risk and healthcare availability.
- World Bank data helped normalize costs using Purchasing Power Parity (PPP) conversion rates.
- UNWTO statistics gave insight into peak tourist seasons, visitor frequency, and seasonal trends.

These sources were merged into a unified dataset containing thousands of city-month records, each with more than 25 attributes.

3.5 Data Preprocessing

Before applying machine learning algorithms, the data had to be cleaned and transformed through several preprocessing steps:

1. **Standardization of Units:** All temperatures were converted to degrees Celsius, and costs were converted to USD using PPP.
2. **Temporal Aggregation:** Daily climate data was averaged to obtain monthly values.
3. **Missing Value Treatment:** Missing climate or cost values were filled using interpolation or median values from nearby regions.
4. **Outlier Removal:** Extreme AQI or crime data were capped to reduce skew.
5. **Encoding of Categorical Fields:** Attributes like continent, health alerts, or seasonality were converted using one-hot encoding or binary flags.
6. **Normalization:** Features like AQI, cost, and temperature were scaled using StandardScaler to prepare for K-Means clustering.

3.6 K-Means Clustering Algorithm

K-Means is an unsupervised algorithm used to group data points into K distinct clusters based on feature similarity. In this project, cities were clustered based on climate (temperature, humidity), environmental conditions (AQI), cost of stay, and health risk indicators.

Why K-Means?

- It helped categorize destinations into meaningful groups such as “hot and dry affordable cities” or “moderate-climate clean cities.”
- Users could then be matched with the most relevant clusters based on their input preferences.

Working of K-Means:

1. Randomly initialize K cluster centroids.
2. Assign each city to the closest centroid based on Euclidean distance.
3. Recalculate the centroids as the mean of assigned points.
4. Repeat the assignment and update steps until centroids stabilize.

```
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Sample input features: [temperature, AQI, cost_index, altitude]
X = [[30, 70, 4.5, 300], [20, 50, 6.0, 1500], [35, 90, 3.2, 200]]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

kmeans = KMeans(n_clusters=2, random_state=42)
kmeans.fit(X_scaled)
print("Cluster labels:", kmeans.labels_)
```

3.7 Nutritional / Cost Prediction Methodology

Some destinations had missing data such as estimated daily cost or AQI. To fill these gaps, a simple prediction method was used.

Steps Involved:

1. Ingredient/Feature Mapping: Using external APIs and statistical summaries, missing values were inferred using data from similar cities.
2. Aggregation of Known Features: Example – if hotel and food costs were known, they were summed to estimate total daily expenses.
3. Regression Modeling: A linear regression model was trained on known data to predict missing values based on features like duration of stay, season, and type of destination.

```
import pandas as pd

from sklearn.linear_model import LinearRegression

df = pd.DataFrame({
    'days': [3, 7, 10],
    'season_code': [0, 1, 1], # 0 = off-peak, 1 = peak
    'hotel_class': [2, 3, 4],
    'total_cost_usd': [300, 950, 1600]
})

X = df[['days', 'season_code', 'hotel_class']]
y = df['total_cost_usd']

model = LinearRegression()
model.fit(X, y)
print("Predicted cost:", model.predict([[5, 1, 3]])[0])
```

3.8 Visualization and Analysis

The insights from clustering and prediction were visualized using scatter plots and heatmaps.

```
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming `df` contains features and assigned cluster labels
sns.scatterplot(x='Temperature', y='AQI', hue='Cluster', data=df)
plt.title('Destination Clusters by Temperature and AQI')
plt.xlabel('Average Monthly Temperature (°C)')
plt.ylabel('Air Quality Index (AQI)')
plt.show()
```

These plots helped validate cluster effectiveness and offered a clear visual interpretation of how destinations were grouped.

By combining clustering, regression, and data visualization techniques, the Smart Travel Advisor system provides intelligent, personalized, and health-conscious travel suggestions that adapt to users' budget, health, and lifestyle requirements.

Chapter 4: Results and Discussion, Performance Analysis

4.1 Destination-Dataset Summary

The consolidated travel dataset used in the Smart Travel Advisor contained **8 127 city-month records**. Each record held core environmental and economic attributes:

- Average monthly temperature (°C)
- Relative humidity (%)
- Air-Quality Index (AQI)
- Cost-of-living index (normalized to USD PPP)
- Altitude (m)
- Monthly rainfall (mm)

A quick statistical scan confirmed substantial diversity:

- Temperatures ranged from **−15 °C** (Arctic towns in January) to **40 °C** (desert cities in July).
- AQI values spanned **15** (mountain hamlets) to **220** (industrial metros in winter).
- Daily cost indices varied almost ten-fold, from backpacker-friendly villages to high-end capitals.

Such wide dispersion justified an **unsupervised clustering approach** to segment destinations into meaningful travel archetypes.

4.2 Optimal Number of Clusters

Using the **Elbow Method** on the Within-Cluster Sum of Squares (WCSS) curve for $k = 1 \dots 10$, the inflection clearly appeared at **$k = 4$** . This choice balanced cluster compactness with human interpretability.

4.3 Cluster Distribution and Centroid Highlights

After training K-Means with $k = 4$, the model produced four distinct destination groups:

1. Cluster A — Clean-Air, Mild-Climate Cities
 - Rough centroid: Temperature ≈ 18 °C, AQI ≈ 35 , cost index ≈ 5.8 (USD PPP scale).
 - Size: 3 212 records (largest group).
2. Cluster B — Budget Warm-Weather Hubs
 - Centroid: Temperature ≈ 30 °C, AQI ≈ 65 , cost index ≈ 3.9 .
 - Size: 2 020 records.
3. Cluster C — High-Altitude, Adventure Zones


- Centroid: Temperature ≈ 12 °C, AQI ≈ 45 , altitude $\approx 1\,800$ m, cost index ≈ 4.7 .
- Size: 1 125 records.

4. Cluster D — Premium Urban Metropolises

- Centroid: Temperature ≈ 25 °C, AQI ≈ 110 , cost index ≈ 8.5 .
- Size: 1 770 records.

4.4 Interpretation of Clusters


- **Cluster A** suits health-sensitive travelers or families wanting moderate temperatures and very low pollution.
- **Cluster B** appeals to budget backpackers who tolerate heat in exchange for affordable living costs.
- **Cluster C** targets trekkers and adventure seekers comfortable with cooler weather and high elevation.
- **Cluster D** attracts business or luxury tourists prepared to pay premium prices and accept denser urban environments



4.5 Visual Validation

Several plots confirmed the quality of grouping:

- **Temperature vs. AQI scatter** colored by cluster showed distinct, minimally overlapping clouds.
- **Cost-of-living histograms** revealed non-intersecting peaks corresponding to each group.
- **3-D plots** (Temperature, AQI, Cost) further illustrated clear separation, validating K-Means effectiveness.



4.6 Performance Analysis

Although clustering is unsupervised and doesn't have traditional accuracy metrics, performance can be discussed in terms of:

Cluster Compactness

- The sharp WCSS drop at $k = 4$ and a mean Silhouette score of 0.51 indicated well-formed, coherent clusters.

Real-world Interpretability


- Each cluster mapped neatly onto recognizable travel types (clean-air retreats, budget tropics, high-altitude adventure, premium metros), making recommendations actionable.

Scalability

- K-Means trained on 8 k+ rows in under two seconds on a standard laptop, confirming feasibility for much larger or near-real-time datasets.

4.7 Clustering driven Use-Case Example

- **Cluster A:** Ideal for asthma patients or parents with infants who require low AQI and mild weather year-round.
- **Cluster B:** Perfect for college students on shoestring budgets seeking warm beaches and street-food culture.
- **Cluster C:** Suits trekkers, mountain cyclists, and digital detox retreats far from pollution.
- **Cluster D:** Fits executives on business trips or luxury travelers interested in fine dining and cosmopolitan nightlife.



Chapter 5: Summary, Conclusions, and Future Scope

Summary

This project demonstrated how unsupervised machine-learning techniques—especially K-Means clustering—can power a Smart Travel Advisor that personalizes destination suggestions according to each traveler’s budget, climate comfort, air-quality needs, and health or activity preferences.

Key steps completed were:

- **Collecting and merging data** from weather APIs, air-quality feeds, cost-of-living portals, and health-safety bulletins to create a rich dataset of more than 8 000 city-month records.
- **Pre-processing** that included unit standardisation ($^{\circ}\text{F} \rightarrow ^{\circ}\text{C}$, local currency $\rightarrow$ USD PPP), missing-value imputation, outlier capping, one-hot encoding of categorical flags, and Z-score normalization of numeric attributes.
- **Implementing K-Means clustering** to segment destinations by shared environmental and economic traits.
- **Selecting the optimal number of clusters** ($k = 4$) using the Elbow Method, which balanced cluster compactness with interpretability.
- **Visualizing and interpreting results** through scatter plots and heatmaps, confirming that clusters aligned with intuitive travel archetypes like “clean-air wellness towns” or “budget tropical hubs.”
- **Generating actionable recommendations** that match user profiles to cluster centroids, thereby offering practical destination short-lists that respect health and budget constraints.

Through this pipeline, the project validated that unsupervised learning can meaningfully organize global travel data and deliver tailored, health-aware travel advice.

Conclusions

- **Effectiveness of Clustering:** K-Means produced clear, well-separated clusters in climate–AQI–cost space, allowing the system to categorize destinations into recognizable travel types.
- **Interpretability and Practical Relevance:** Cluster centroids mapped cleanly to terms travelers understand (e.g., *budget warm-weather*, *premium urban*, *high-altitude adventure*), making recommendations easy to explain and act upon.
- **Data-Driven Travel Planning:** By basing suggestions on real temperature, air-quality, and cost data, the Advisor empowers users to make informed decisions rather than relying on hearsay or generic popularity scores.
- **Performance and Scalability:** The workflow handled thousands of records in seconds on modest hardware, showing that it can scale to much larger datasets or near-real-time updates without heavy computational overhead.

Overall, the Smart Travel Advisor proves that unsupervised learning is a viable backbone for modern, personalized travel-recommendation engines.

Future Scope

1. Deeper Personalization

Add user accounts that store medical conditions, altitude tolerance, and preferred activity styles.

Learn from past clicks and bookings to refine future suggestions dynamically.

2. Advanced and Adaptive Clustering

Experiment with DBSCAN, Gaussian Mixture Models, or hierarchical clustering to capture non-spherical or overlapping destination groups.

Combine clustering with dimensionality-reduction techniques (e.g., UMAP) to visualize more complex feature sets.

3. Real-Time Recommendation Service

Deploy the pipeline as a web or mobile app delivering instant recommendations based on live user inputs.

Stream continuous data feeds for weather alerts, pollution spikes, and currency fluctuations.

4. Expanded Data Integration

Ingest live flight prices, hotel availability, and local event calendars via APIs to enhance cost and experience predictions.

Include crime indices, internet-speed maps, and visa-requirement databases for richer decision factors.

5. Hybrid Recommendation Models

Combine cluster-based filtering with collaborative or content-based methods to incorporate community ratings and behavioral trends.

Explore deep-learning approaches—such as graph neural networks—to model relationships between users, destinations, and travel seasons.

6. Health and Sustainability Linkages

Sync with wearable devices and health apps to adjust recommendations based on exertion levels or respiratory alerts.

Provide carbon-footprint estimates and eco-friendly travel alternatives to support sustainable tourism goals.

References

Journals

1. Jain, A. K. (2010). "Data clustering: 50 years beyond K-means." *Pattern Recognition Letters*, **31**(8), 651–666. <https://doi.org/10.1016/j.patrec.2009.09.011>
2. Aggarwal, C. C., & Reddy, C. K. (2013). "Data Clustering: Algorithms and Applications." *CRC Press*, **1st Edition**, pp. 45–67.
3. Huang, Z., & Ling, C. X. (2005). "Using A Machine Learning Approach to Predict Diet Recommendation for Diabetics." *International Journal of Computer Science and Network Security*, **5**(8), 248–255.
4. Liu, S., Cao, J., & Yu, Z. (2020). "A Personalized Diet Recommendation System Based on Ontology and Machine Learning." *IEEE Access*, **8**, 56995–57004.

Books

1. Han, J., Kamber, M., & Pei, J. (2011). *Data Mining: Concepts and Techniques*, 3rd Edition, Edited by Morgan Kaufmann Publishers, Elsevier, pp. 443–454.
2. Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*, Springer, New York, pp. 289–296.
3. Mitchell, T. M. (1997). *Machine Learning*, McGraw-Hill, International Edition, pp. 52–54.

Internet Articles

1. Scikit-Learn Documentation. "K-Means Clustering." *Scikit-Learn Developers*. <https://scikit-learn.org/stable/modules/clustering.html#k-means>
2. USDA FoodData Central. "Nutrient Values for Common Foods." *U.S. Department of Agriculture*. <https://fdc.nal.usda.gov>
3. Edamam Nutrition API. "Meal Planning with Nutritional Analysis." *Edamam Inc*. <https://developer.edamam.com>
4. Towards Data Science. "How to Use KMeans Clustering for Data Segmentation." *Medium.com*. <https://towardsdatascience.com/k-means-clustering-algorithm-applications-evaluation-methods-and-drawbacks-aa03e644b48a>
5. Harvard T.H. Chan School of Public Health. "The Nutrition Source." <https://www.hsph.harvard.edu/nutritionsource/>

APPENDIX

```
import pandas as pd

import numpy as np

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.ensemble import RandomForestClassifier

from sklearn.cluster import KMeans

from sklearn.preprocessing import StandardScaler, LabelEncoder

from sklearn.metrics import mean_squared_error, classification_report


trip_df = pd.read_csv('travel_trip_data.csv')

cost_df = pd.read_csv('cost_data.csv')

weather_df = pd.read_csv('weather_data.csv')

health_df = pd.read_csv('health_data.csv')

print(trip_df)

print(cost_df)

print(weather_df)

print(health_df)


print("Trip Data Info:")

print(trip_df.info())

print(trip_df.describe())


print("\nCost Data Info:")
```

```
print(cost_df.info())
print(cost_df.describe())

print("\nWeather Data Info:")
print(weather_df.info())
print(weather_df.describe())

print("\nHealth Data Info:")
print(health_df.info())
print(health_df.describe())

# Display info and summary stats
print("Trip Data Info:")
trip_df.info()
print("\nTrip Data Description:")
print(trip_df.describe(include='all'))

print("\nCost Data Info:")
cost_df.info()

print("\nWeather Data Info:")
weather_df.info()

print("\nHealth Data Info:")
health_df.info()
```

```

# Visualizations for Trip Data

sns.histplot(trip_df['Accommodation cost'].dropna(), kde=True)

plt.title('Distribution of Accommodation Cost')

plt.show()


sns.histplot(trip_df['Transportation cost'].dropna(), kde=True)

plt.title('Distribution of Transportation Cost')

plt.show()

trip_df['Accommodation cost'] = pd.to_numeric(trip_df['Accommodation cost'],
errors='coerce')

trip_df['Transportation cost'] = pd.to_numeric(trip_df['Transportation cost'], errors='coerce')

trip_df['Start date'] = pd.to_datetime(trip_df['Start date'], errors='coerce')

trip_df['End date'] = pd.to_datetime(trip_df['End date'], errors='coerce')

trip_df.dropna(subset=['Destination', 'Start date', 'End date', 'Accommodation cost',
'Transportation cost'], inplace=True)

trip_df['Total Trip Cost'] = trip_df['Accommodation cost'] + trip_df['Transportation cost']

trip_df['Country'] = trip_df['Destination'].apply(lambda x: x.split(',')[1].strip() if isinstance(x,
str) and ',' in x else np.nan)

print(trip_df)


sns.histplot(trip_df['Total Trip Cost'].dropna(), kde=True)

plt.title('Distribution of Total Trip Cost')

plt.show()


sns.boxplot(x=trip_df['Country'], y=trip_df['Total Trip Cost'])

plt.xticks(rotation=45)

plt.title('Trip Cost by Country')

plt.show()

```

```

sns.pairplot(trip_df[['Duration (days)', 'Traveler age', 'Total Trip Cost']].dropna())

plt.suptitle('Pairplot of Trip Features', y=1.02)

plt.show()

weather_df_cleaned = weather_df.dropna(subset=['temperature', 'humidity', 'Country', 'City'])
weather_cost = pd.merge(weather_df_cleaned, cost_df, on='Country', how='left')
trip_merged = pd.merge(trip_df, weather_cost, on='Country', how='left')
print(trip_merged)


features = ['Duration (days)', 'Traveler age']
trip_cost_data = trip_merged.dropna(subset=features + ['Total Trip Cost'])
X = trip_cost_data[features]
y = trip_cost_data['Total Trip Cost']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print(X_train)
print(X_test)
print(y_train)
print(y_test)


model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Trip Cost Prediction RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))


health_df = health_df.copy()
health_df = health_df[health_df['Cases'].apply(lambda x: isinstance(x, (int, float)) or
(isinstance(x, str) and x.replace('.', '', 1).isdigit()))]

```



```

health_df['Cases'] = health_df['Cases'].astype(float)

health_df.dropna(subset=['Temp', 'preci', 'LAI', 'Cases'], inplace=True)

Xh = health_df[['Temp', 'preci', 'LAI']]

yh = health_df['Cases']

Xh_train, Xh_test, yh_train, yh_test = train_test_split(Xh, yh, test_size=0.2,
random_state=42)

health_model = LinearRegression()

health_model.fit(Xh_train, yh_train)

yh_pred = health_model.predict(Xh_test)


print("Disease Case Prediction RMSE:", np.sqrt(mean_squared_error(yh_test, yh_pred)))


trip_merged['Safe'] = trip_merged['temperature'].apply(lambda x: 1 if 15 <= x <= 30 else 0)

clf_data = trip_merged.dropna(subset=features + ['Safe'])

Xc = clf_data[features]

yc = clf_data['Safe']

Xc_train, Xc_test, yc_train, yc_test = train_test_split(Xc, yc, test_size=0.2,
random_state=42)

print(Xc_test)

print(yc_train)

print(yc_test)

print(Xc_train)


clf = RandomForestClassifier(random_state=42)

clf.fit(Xc_train, yc_train)

yc_pred = clf.predict(Xc_test)

print("\nTravel Suitability Classification Report:\n")

print(classification_report(yc_test, yc_pred))

```

```

cluster_data = trip_merged[['temperature', 'humidity', 'Cost of Living Index']].dropna()

scaler = StandardScaler()

scaled = scaler.fit_transform(cluster_data)

kmeans = KMeans(n_clusters=3, random_state=42)

clusters = kmeans.fit_predict(scaled)

cluster_data['Cluster'] = clusters

print(cluster_data)

from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5],
}

rf_clf = RandomForestClassifier(random_state=42)

grid_search = GridSearchCV(rf_clf, param_grid, cv=5, scoring='accuracy', n_jobs=-1)

grid_search.fit(Xc_train, yc_train)

print("Best parameters from GridSearchCV:", grid_search.best_params_)

print("Best cross-validation score:", grid_search.best_score_)

best_rf_clf = grid_search.best_estimator_

y_pred_best = best_rf_clf.predict(Xc_test)

print("Classification report with tuned RF model:")

print(classification_report(yc_test, y_pred_best))

```

```

from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import RandomForestClassifier

from sklearn.naive_bayes import GaussianNB

from sklearn.neural_network import MLPClassifier

from sklearn.metrics import classification_report, accuracy_score

def evaluate_classifiers(X_train, y_train, X_test, y_test):

    classifiers = {

        'Random Forest': RandomForestClassifier(random_state=42),

        'Decision Tree': DecisionTreeClassifier(random_state=42),

        'Gaussian NB': GaussianNB(),

        'MLP': MLPClassifier(hidden_layer_sizes=(50,), max_iter=500, random_state=42)

    }

    results = []

    for name, clf in classifiers.items():

        clf.fit(X_train, y_train)

        y_pred = clf.predict(X_test)

        acc = accuracy_score(y_test, y_pred)

        results.append((name, acc))

        print(f'{name} Accuracy: {acc:.4f}')

        print(classification_report(y_test, y_pred))

    return results

results = evaluate_classifiers(Xc_train, yc_train, Xc_test, yc_test)

```

```

names, scores = zip(*results)

sns.barplot(x=scores, y=names)

plt.xlabel("Accuracy")

plt.title("Classifier Performance Comparison")

plt.show()


sns.scatterplot(data=cluster_data, x='temperature', y='Cost of Living Index', hue='Cluster',
palette='Set1')

plt.title('Travel Destination Clusters')

plt.xlabel('Temperature')

plt.ylabel('Cost of Living Index')

plt.show()

# Predict Trip Cost

input_trip = pd.DataFrame({'Duration (days)': [2], 'Traveler age': [19]})

predicted_cost = model.predict(input_trip)

print("\nPredicted Trip Cost for 10 days & age 30:", predicted_cost[0])


# Predict Disease Cases

input_health = pd.DataFrame({'Temp': [27.0], 'preci': [120.0], 'LAI': [2.5]})

predicted_cases = health_model.predict(input_health)

print("Predicted Disease Cases for given weather conditions:", predicted_cases[0])


# Check Travel Suitability

input_suitability = pd.DataFrame({'Duration (days)': [7], 'Traveler age': [22]})

is_suitable = clf.predict(input_suitability)

print("Is the trip suitable?", "Yes" if is_suitable[0] == 1 else "No")

```

```

# Predict Destination Cluster

input_cluster = pd.DataFrame({'temperature': [39], 'humidity': [91], 'Cost of Living Index':
[46]})

input_cluster_scaled = scaler.transform(input_cluster)

predicted_cluster = kmeans.predict(input_cluster_scaled)

print("Predicted Destination Cluster:", predicted_cluster[0])

def classify_temp(temp):
    if temp < 15:
        return 0 # Cold
    elif 15 <= temp <= 30:
        return 1 # Moderate (Ideal)
    else:
        return 2 # Hot

trip_merged['Safe_Level'] = trip_merged['temperature'].apply(classify_temp)

clf_data = trip_merged.dropna(subset=features + ['Safe_Level'])

Xc = clf_data[features]
yc = clf_data['Safe_Level']

Xc_train, Xc_test, yc_train, yc_test = train_test_split(Xc, yc, test_size=0.2,
random_state=42)

clf = RandomForestClassifier(random_state=42)

clf.fit(Xc_train, yc_train)

yc_pred = clf.predict(Xc_test)

# 3x3 confusion matrix

cm = confusion_matrix(yc_test, yc_pred, labels=[0, 1, 2])

```

```
ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Cold", "Moderate",  
"Hot"]).plot()
```

```
plt.title("3x3 Confusion Matrix for Travel Suitability")
```

```
plt.show()
```

```
plt.scatter(yh_test, yh_pred, alpha=0.5)
```

```
plt.xlabel("Actual Disease Cases")
```

```
plt.ylabel("Predicted Disease Cases")
```

```
plt.title("Health Risk Prediction")
```

```
plt.show()
```

```
plt.scatter(y_test, y_pred, alpha=0.5)
```

```
plt.xlabel("Actual Total Trip Cost")
```

```
plt.ylabel("Predicted Total Trip Cost")
```

```
plt.title("Trip Cost Prediction")
```

```
plt.show()
```